



Which deploy spiked your LLM bill?

An autonomous agent for LLM cost attribution.

THE PROBLEM

Your LLM bill doubled. Nobody knows why.


- › Spend dashboards show **that** cost went up. They never show **what shipped** that made it go up.
- › A model swap in one pull request, a retry loop in another, and the bill climbs quietly for days.
- › Someone burns an afternoon cross-referencing Git history against billing exports, by hand.

Observability tools stop at the symptom

TOOL	WHAT IT GIVES YOU	WHAT IT MISSES
Helicone	Hosted observability and caching	No deploy correlation
Langfuse	LLM trace observability	Trace-focused, not cost-causal
LiteLLM	Open-source proxy	No anomaly detection or deploy join
Ilmtrace	Self-hosted proxy and causal agent	Names the deploy that caused it

The sharp claim: deploy to LLM-spend causality, with zero hosted SaaS dependency.

It joins the spike to the deploy that caused it



Anomaly detection, time-window correlation, and lineage scoring, run by an agent instead of an engineer.

A proxy, a ledger, and a Gemini agent



Point your code at Ilmtrace instead of the provider URL. Nothing else changes.

Three tool calls to the root cause

```
anomaly: key=prod-frontend  $19.20 vs $4.48 baseline  28.0σ
```

```
[tool] query_model_distribution  prod-frontend
```

```
[tool] get_deploys_in_window    pivot 2026-05-18T09:40Z
```

```
[tool] diff_prompt_model_mix    before 91% haiku / after 89% sonnet
```

Attribution: PR #2 "switch summary endpoint to claude-sonnet"
shifted the summary prompt from claude-haiku to a costlier model.

Confidence: 0.95

It does not wait to be asked

Background watcher

Scans the ledger on a loop. A new spike triggers an investigation with no human in the loop.

It takes action

Opens a GitHub issue, tightens a rate limit, or writes a corrected file and opens a fix pull request.

Real GitHub actions

The agent reads PR diffs and source files, then ships a real branch and PR to your repo.

Every action logged

The dashboard shows each scan, attribution, and remediation as it happens, live.

Vision Import: drop in **any** **billing screenshot**

- › No ledger yet? Upload a screenshot of any billing dashboard, Anthropic, OpenAI, or a Grafana panel.
- › **Gemini multimodal** reads the daily spend straight off the image and finds the spike.
- › The agent then investigates your connected GitHub repo and attributes the cause, end to end.

The same causal engine, reachable from a picture.

One Go binary. Zero hosted SaaS.

Go

Proxy, ledger, anomaly detection, and web server on the standard library.

Gemini

google.golang.org/genai SDK, multi-turn tool-calling agent loop.

SQLite

Pure-Go modernnc.org/sqlite, no CGo. The whole ledger is one file.

Google Cloud Run

Single container, deployed live and serving real traffic right now.

TRY IT

10 / 10

Live right now

LIVE DEMO llmtrace-681081536857.asia-south1.run.app

SOURCE github.com/Yatsuiii/llmtrace

STACK Go · Gemini · SQLite · Cloud Run

Part of the trace family: operational tooling that joins telemetry to the change that caused it, across cloud and AI ops.